



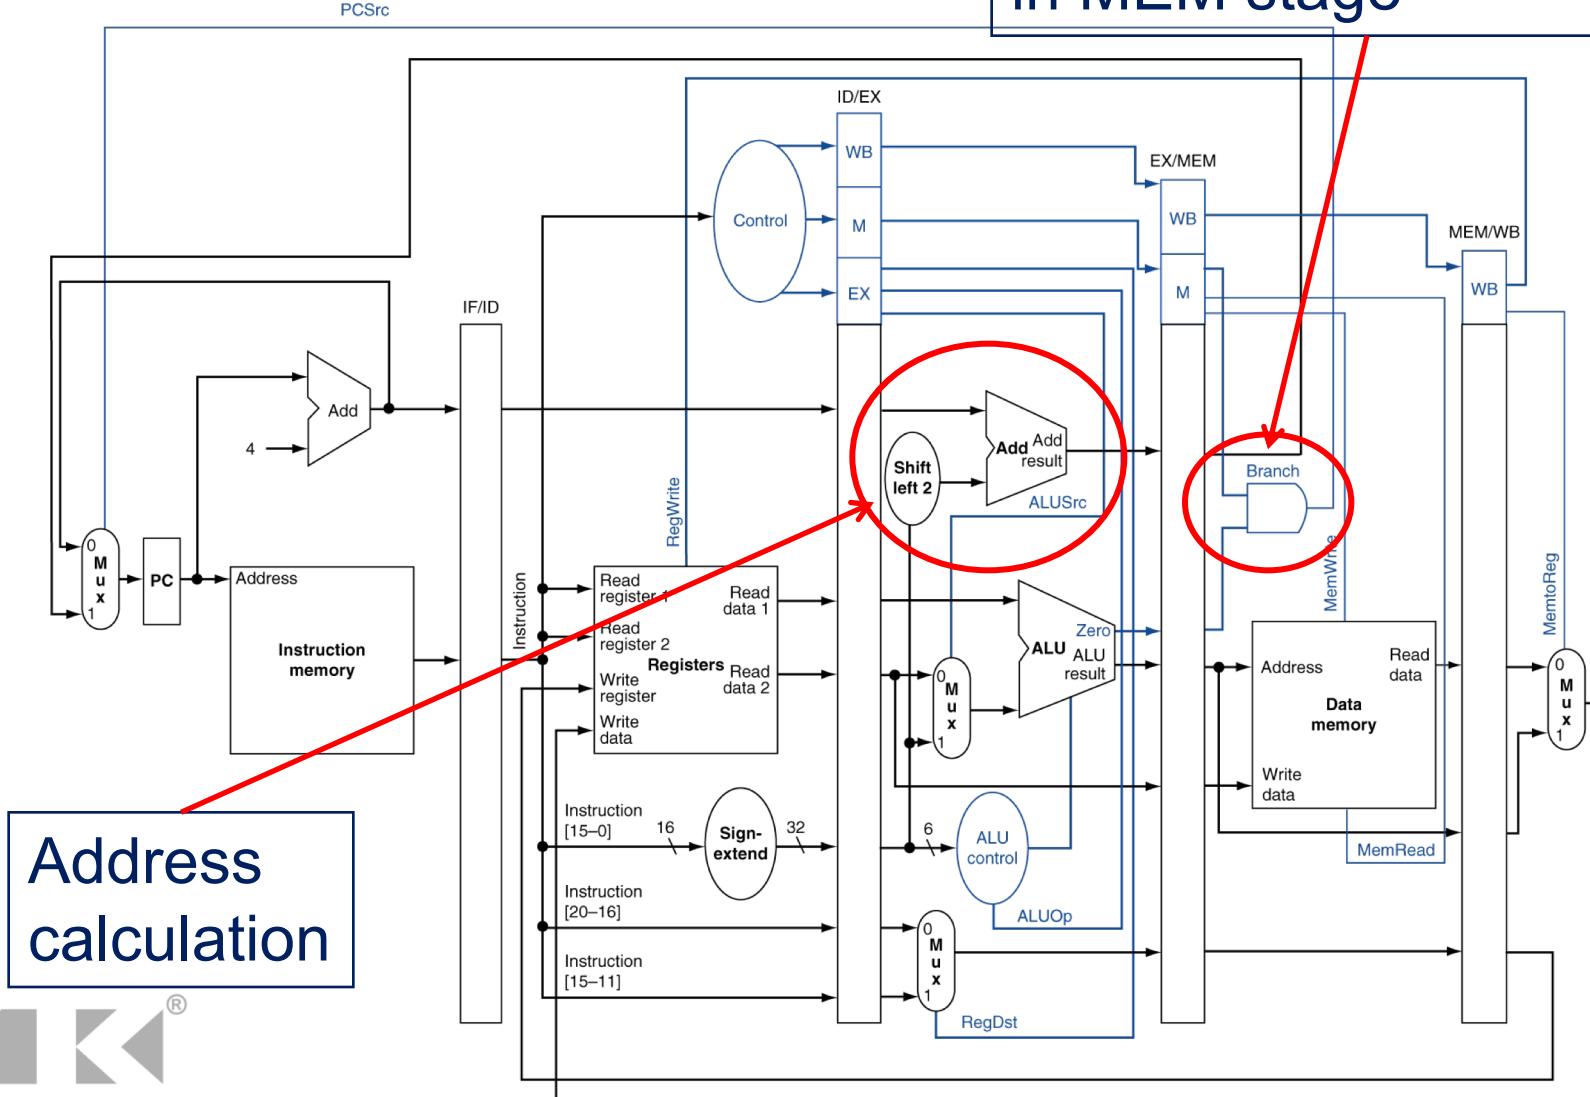
Topic 9

Control Hazards

Branch Hazards

Current implementation

Determination for branch
in MEM stage



Control Hazards

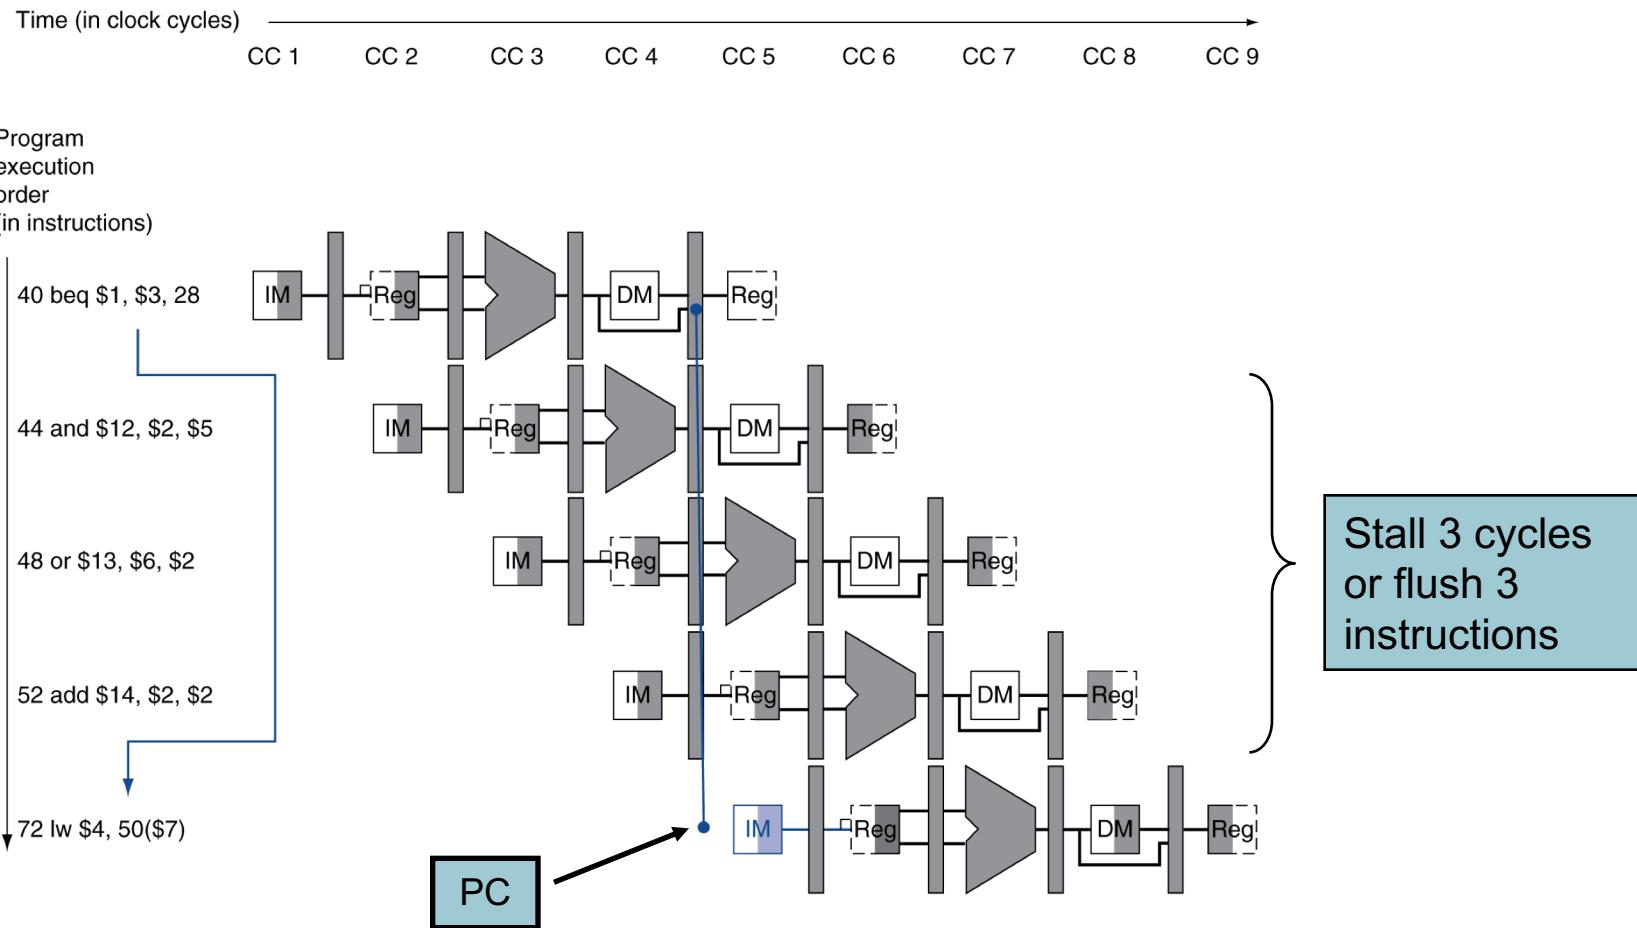
- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction (Branch Hazard) because
 - When Branch instruction is still in the ID stage, target instruction is needed in the IF stage

Branch Hazard Resolutions

- *Stall on branch*
- Always assume branch not taken or taken
- Branch prediction
- Delayed Branch

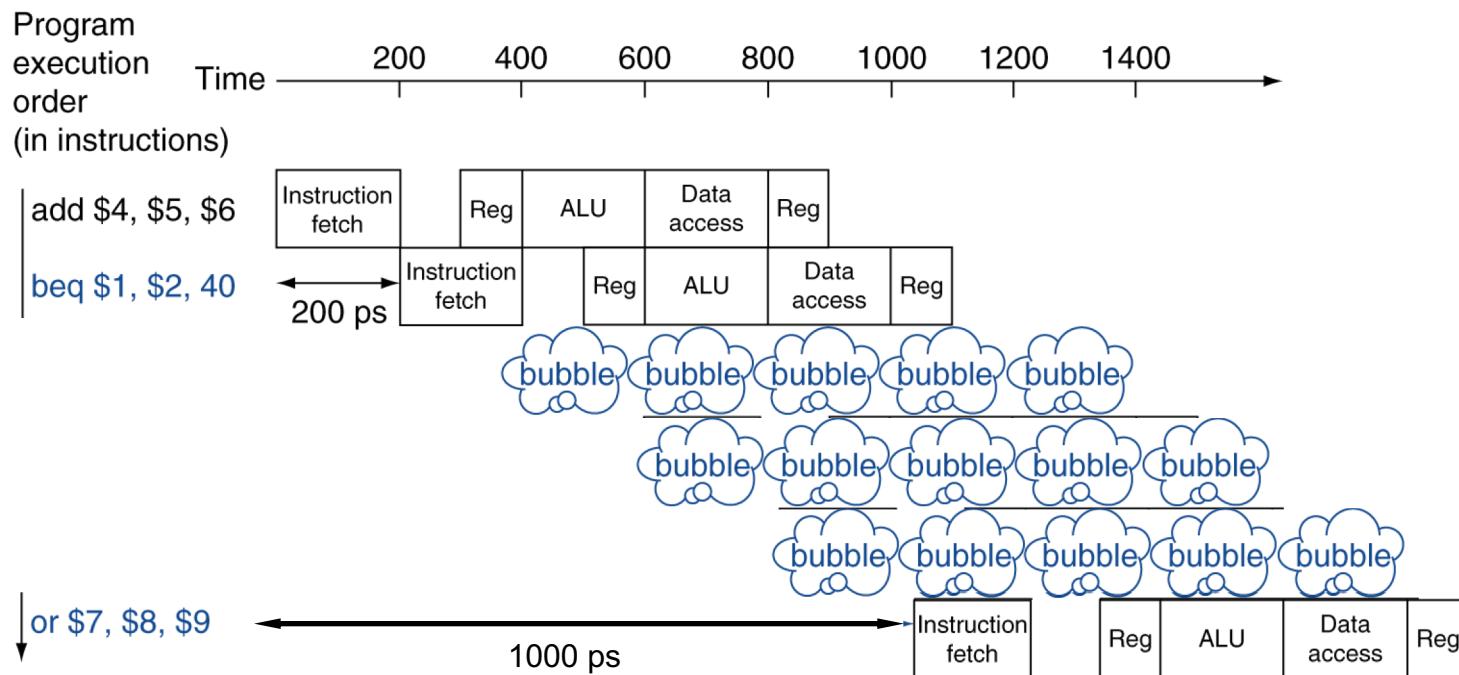
Branch Hazards

- If branch outcome determined in MEM
 - Stall branch or take penalty of wrong branch



Stall on Branch

- Wait until branch outcome determined before fetching next instruction



Branch Hazard Resolutions

- Stall on branch
- *Always assume branch not taken or taken*
- Branch prediction
- Delayed Branch

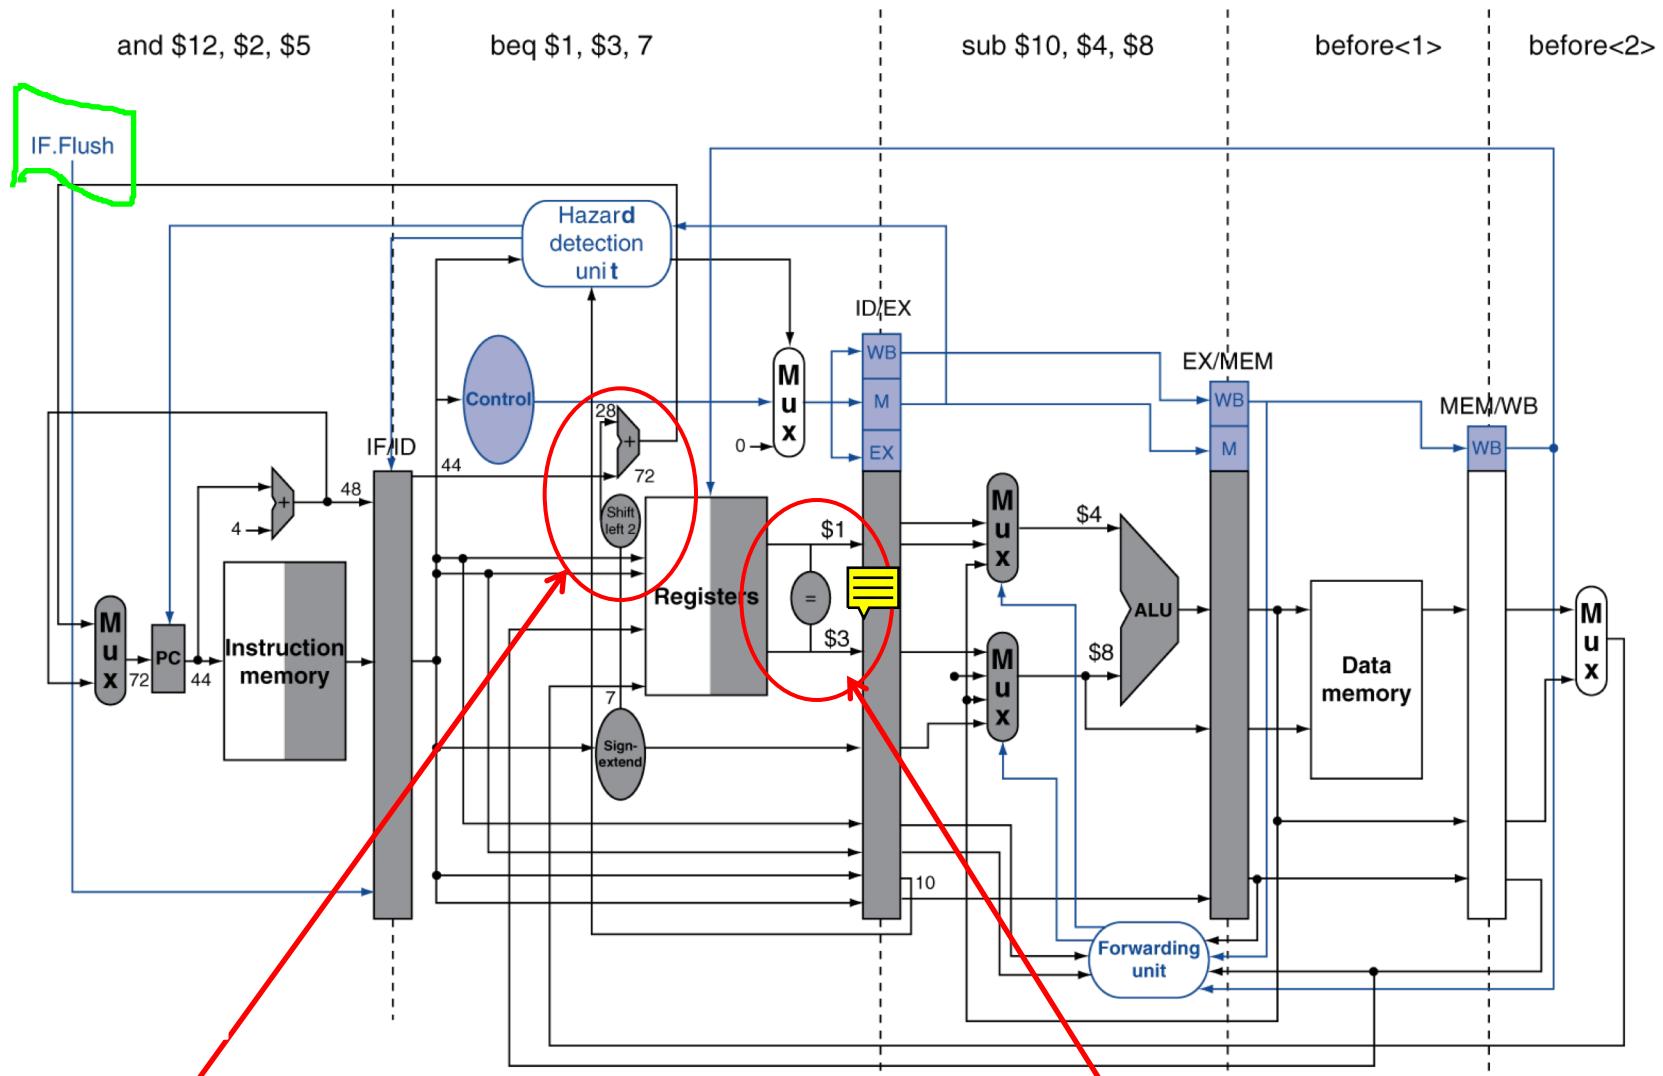
Assume Branch Not Taken

- If we are right, lucky us!
- If we are wrong, penalty will be to flush some (up to 3) instructions
- Penalty may be reduced
 - By making earlier decision on branch or not
 - Need to compare registers and compute target early in the pipeline

Reducing Branch Penalty

- Move hardware for determining PC to *ID stage* including
 - Target address calculation
 - Register comparator
- For branch not taken, correct assumption, no penalty
- For taken branch, penalty reduced
 - 1 instruction to be flushed

Example: Branch Taken



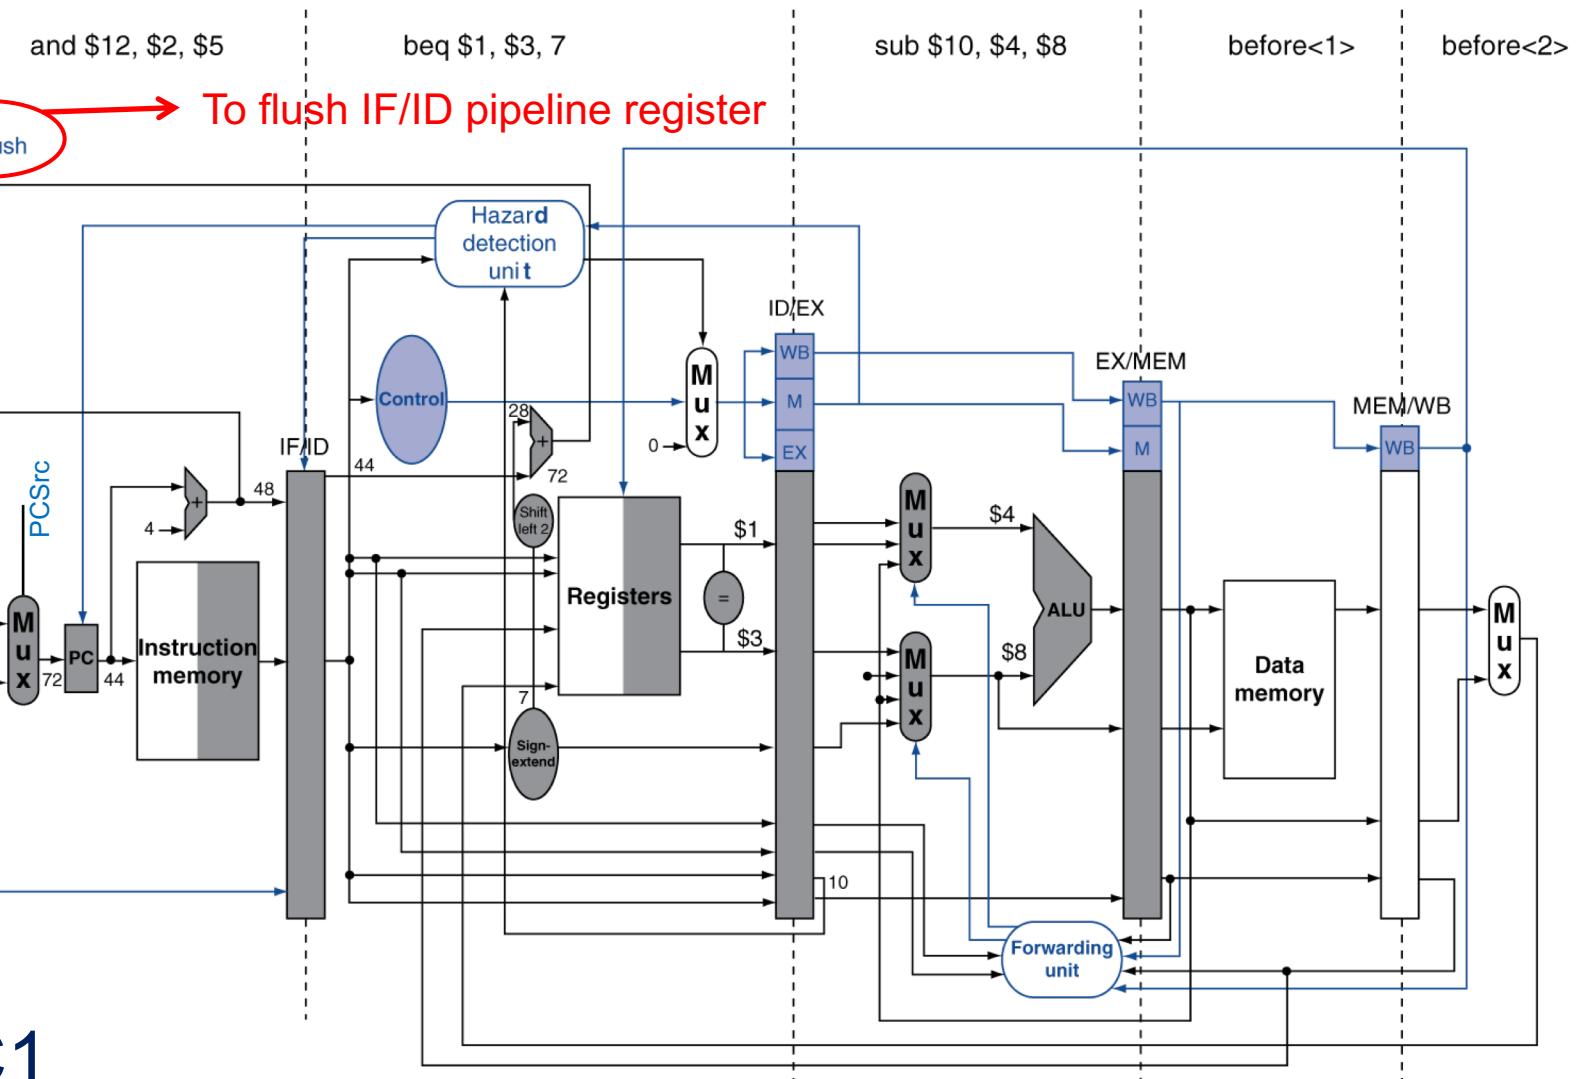
Address calculation

Register comparison

Flush an Instruction

- Flush: To discard the wrong instruction in pipeline, equivalent to neutralize all operations
 - Clear IF/ID pipeline register, by a new control signal **IF.Flush**
 - Flushes the instruction in IF stage

IF.Flush (synchronous)	Branch	“=” Output	PCSrc	PCin	PCout	IF/ID Register
1	1	1	0	72	44	beq & PC+4



CC1

IF.Flush (synchronous)

Branch

“=” Output

PCSrc

IF/ID Register

0

0

1

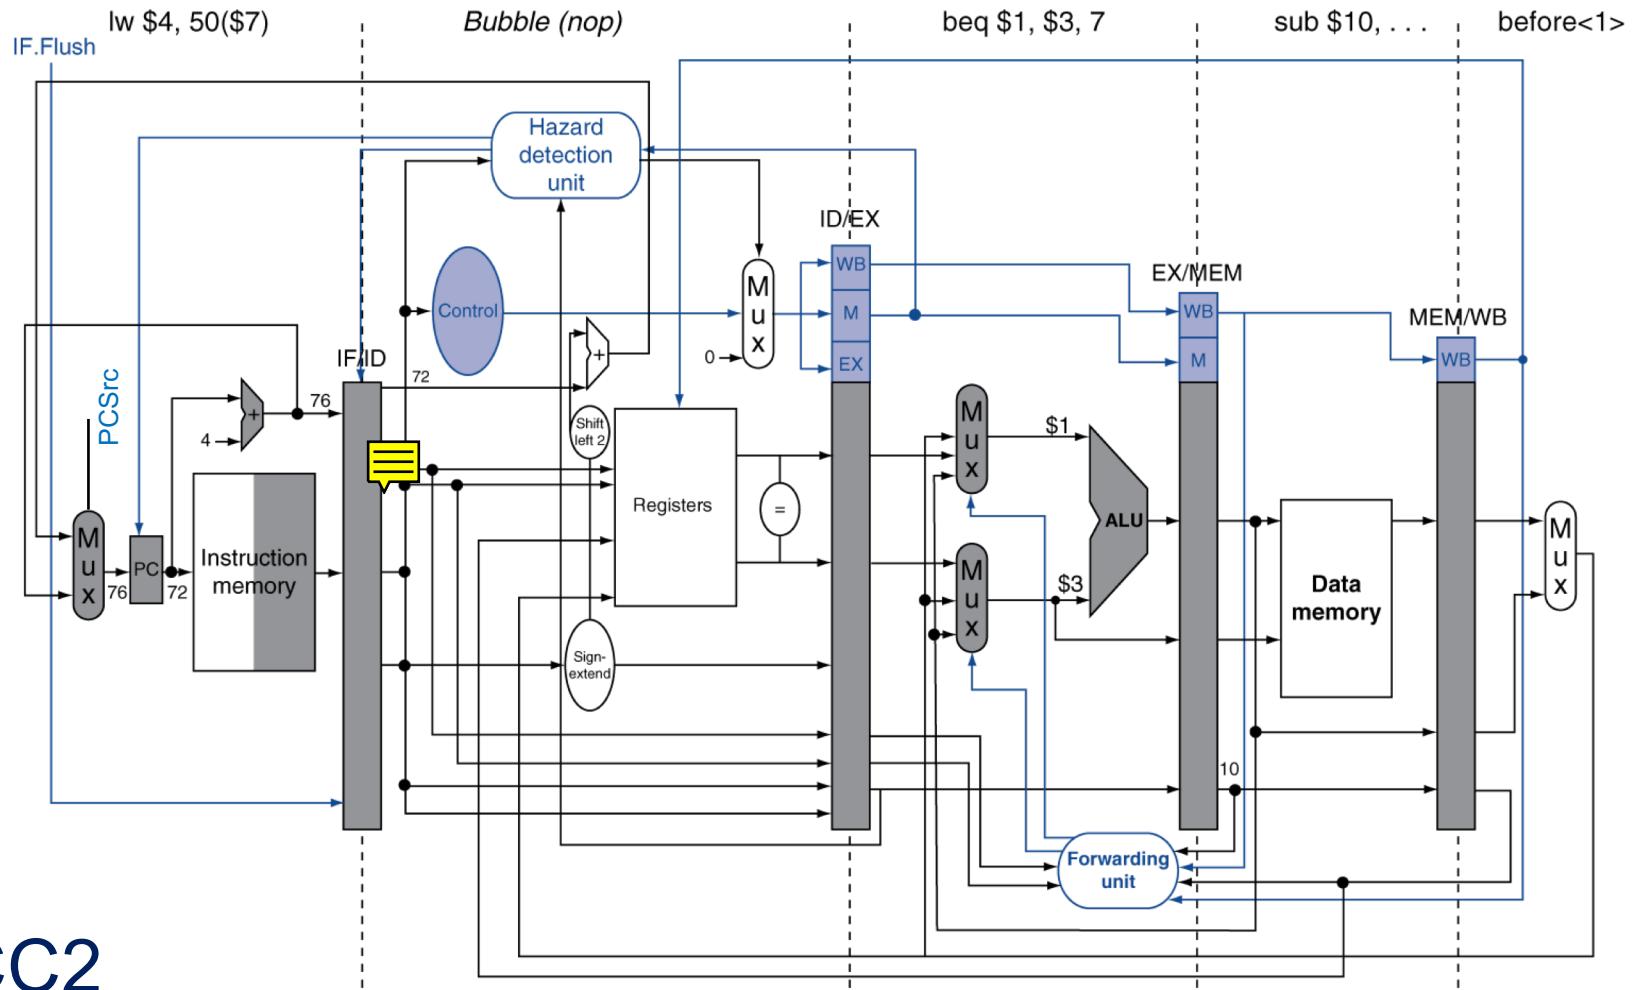
1

0

Branch target instruction

“and” instruction flushed

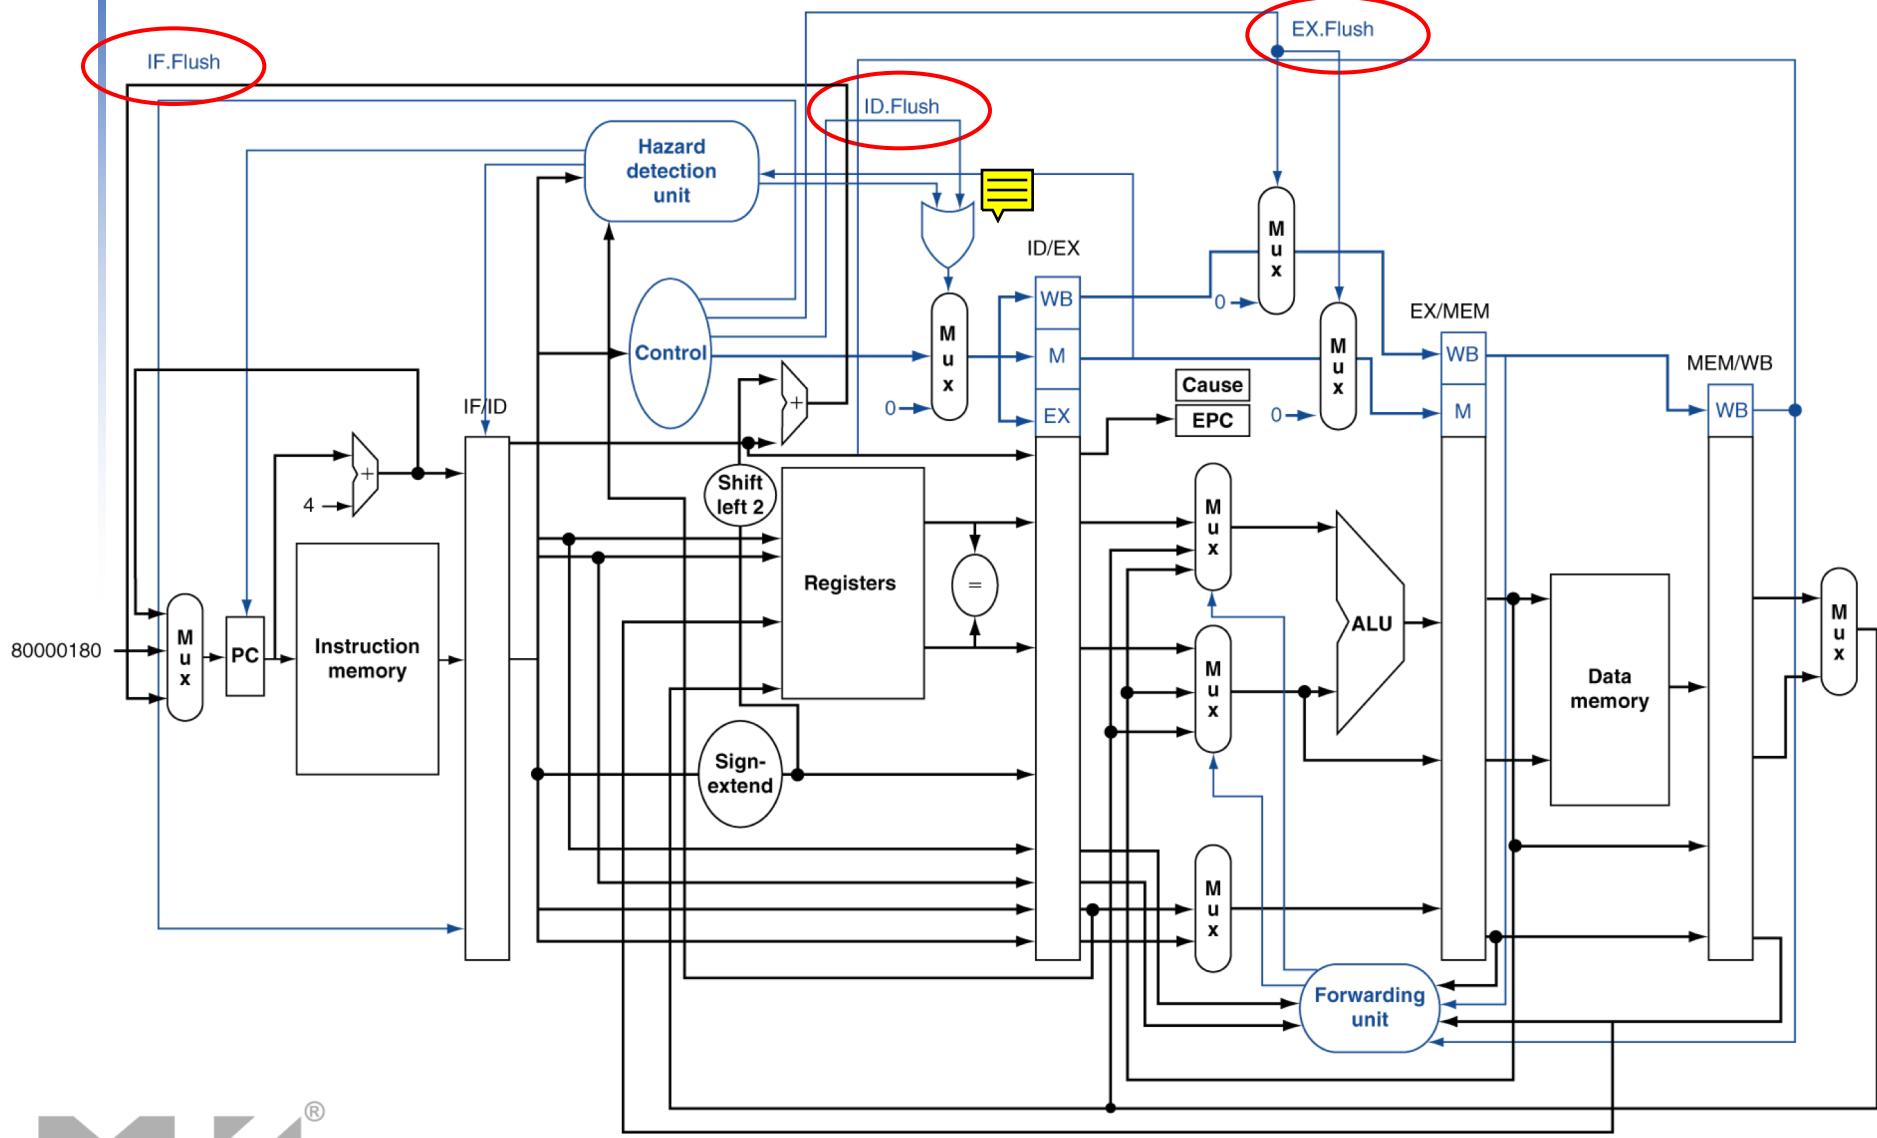
Branch condition is true



Flush Multiple Instructions

- What's in IF stage is previous 1 instruction, how to flush previous 2 or 3 instructions in ID and EX stages?
 - Send 0 control signals through the pipeline
 - Flushes the instruction in ID stage
 - Flushes instructions in following stage

Pipeline with Extra Controls



IF.Flush**ID.Flush****EX.Flush**

1

1

1

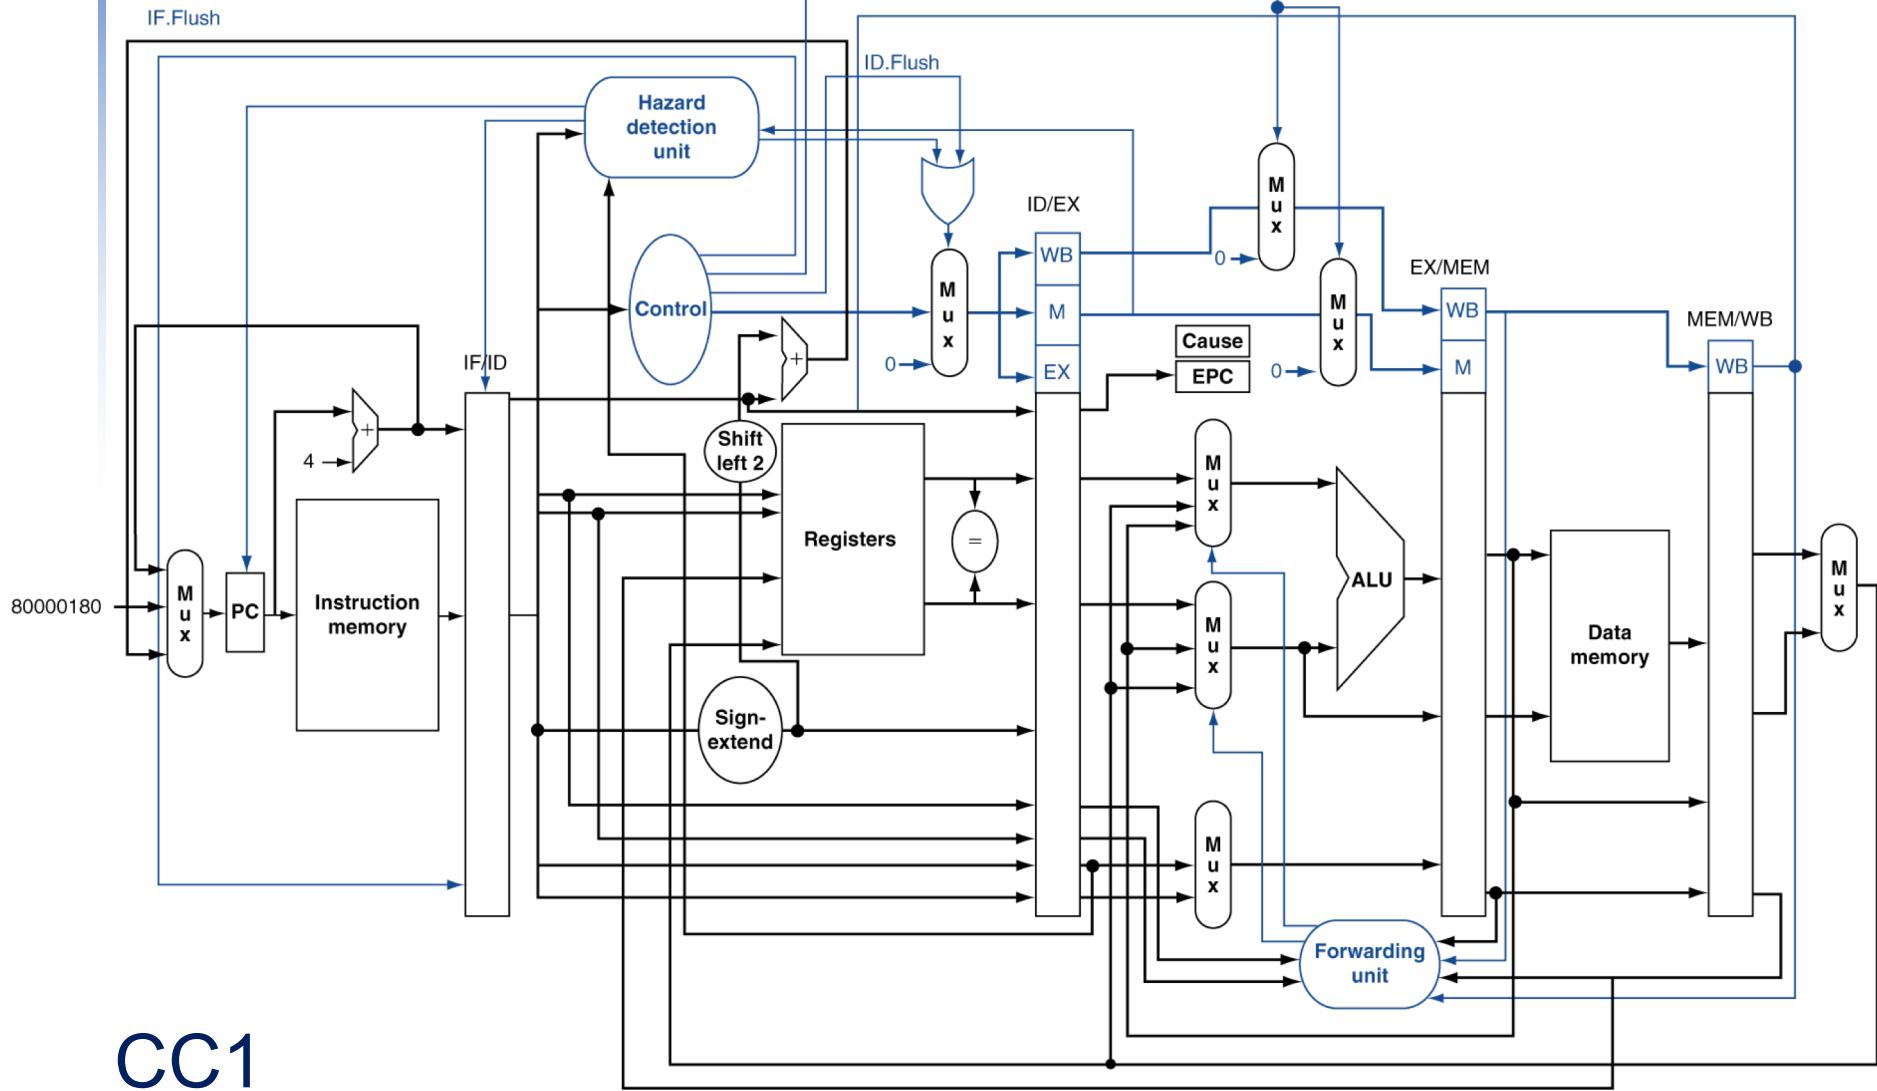
add

or

and

beq

add



IF.Flush**ID.Flush****EX.Flush**

0

0

0

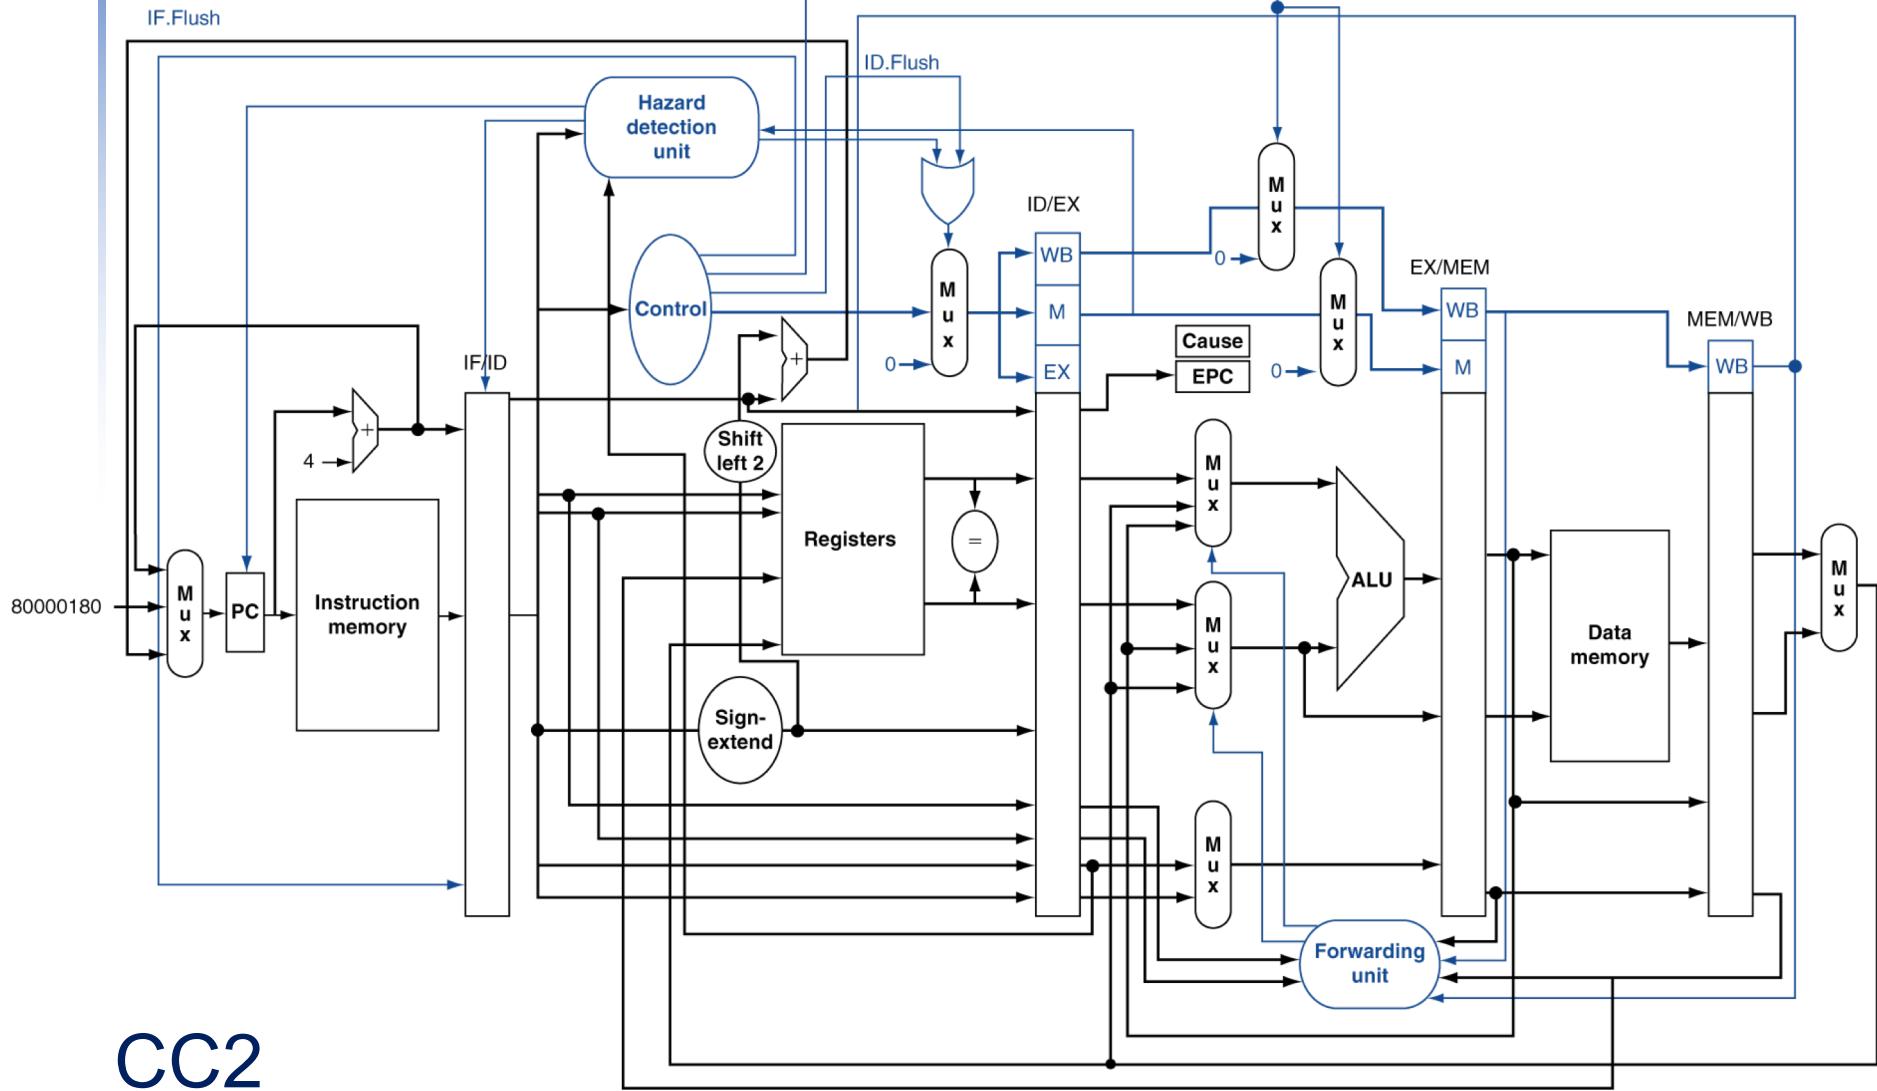
addi

nop

nop

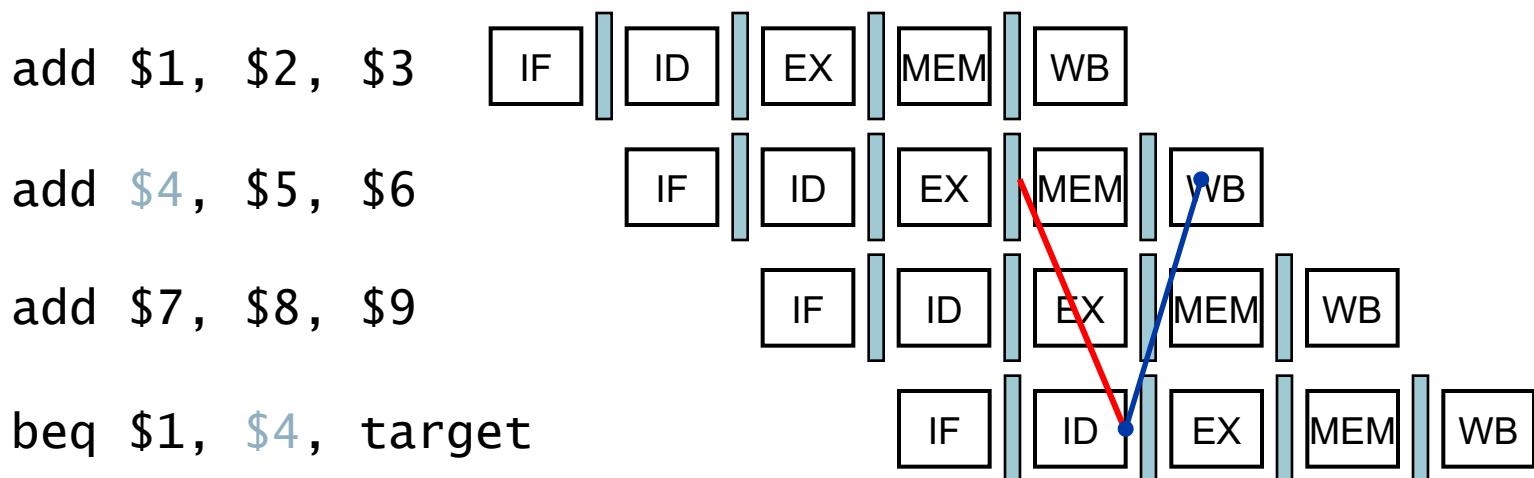
nop

beq



Data Hazards for Branches

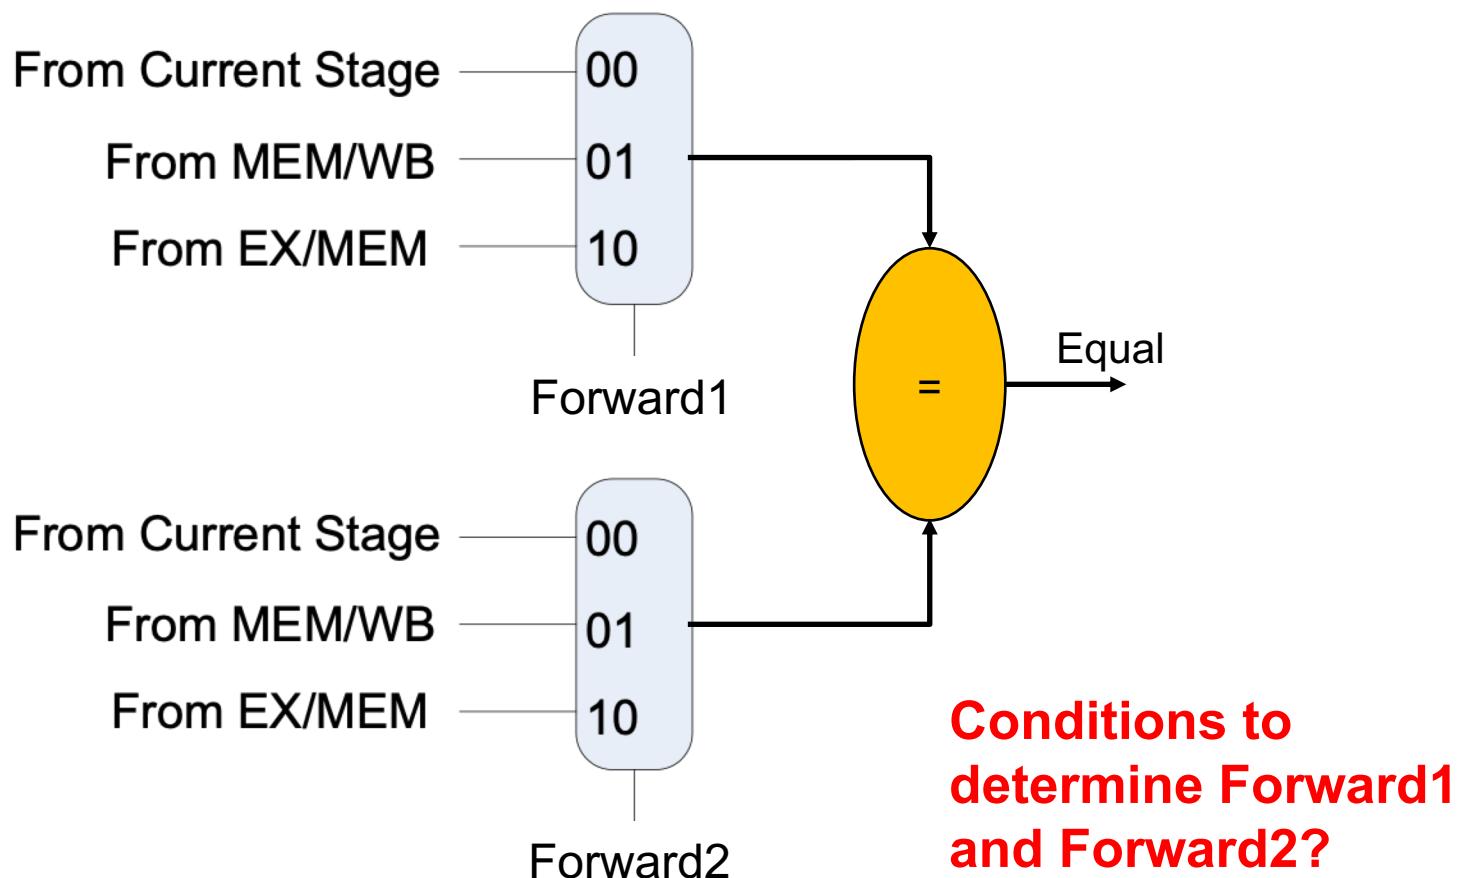
- Changing datapath structure causes more possible data hazards
- If a register for comparison is a destination of 2nd preceding ALU instruction



- Can resolve using new forwarding paths
- Is \$1 a data hazard?

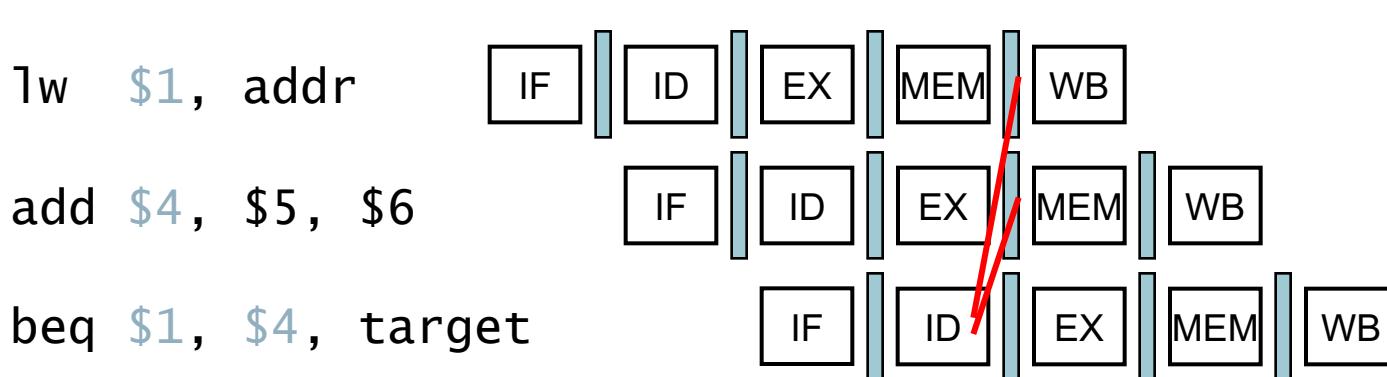
Forwarding Paths

- Forwarding paths are created between stage pipeline registers and comparator inputs



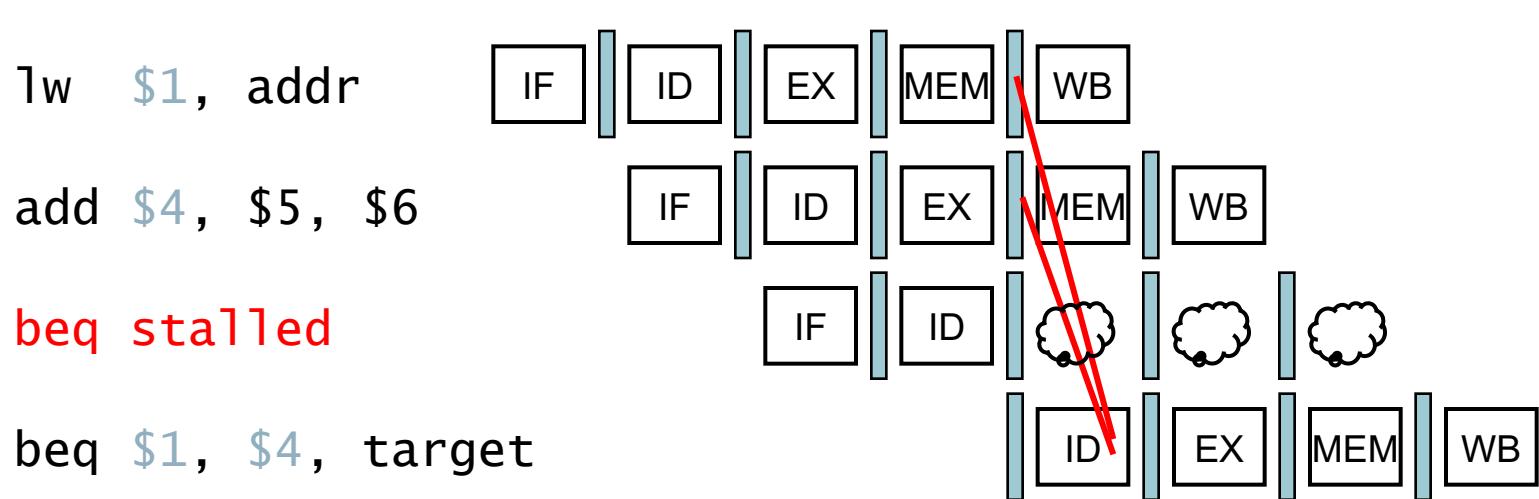
Data Hazards for Branches

- If a comparison register is a destination of *immediately preceding* ALU instruction or 2^{nd} preceding load instruction



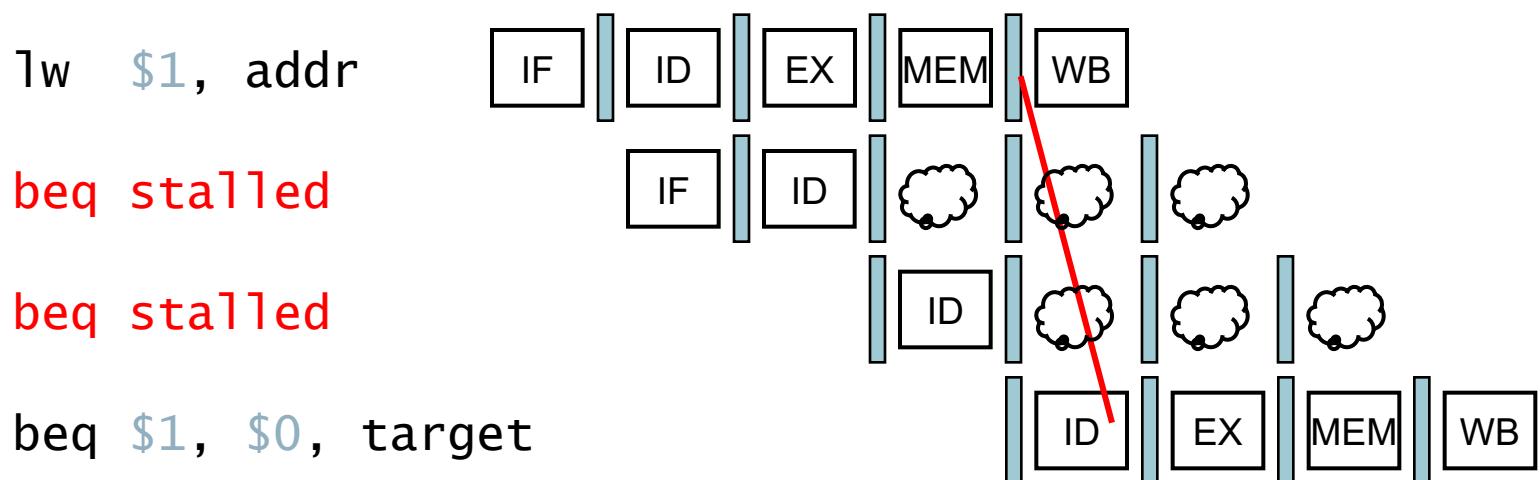
Data Hazards for Branches

- Need 1 stall cycle even with forwarding



Data Hazards for Branches

- If a comparison register is a destination of *immediately preceding* load instruction
 - Need 2 stall cycles



Branch Hazard Resolutions

- Stall on branch
- Always assume branch not taken or taken
- *Branch prediction (instead of assumption)*
- Delayed Branch

Branch Prediction

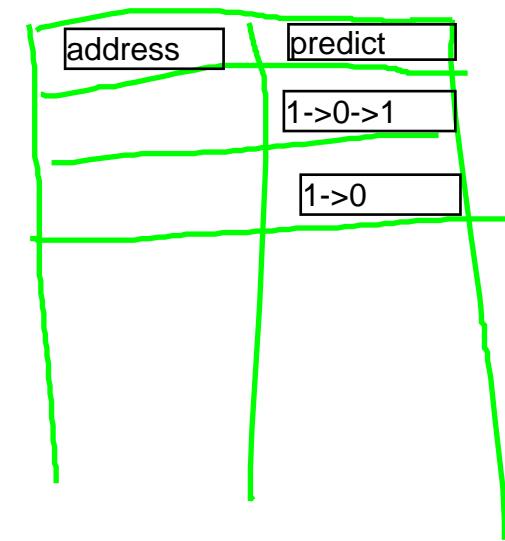
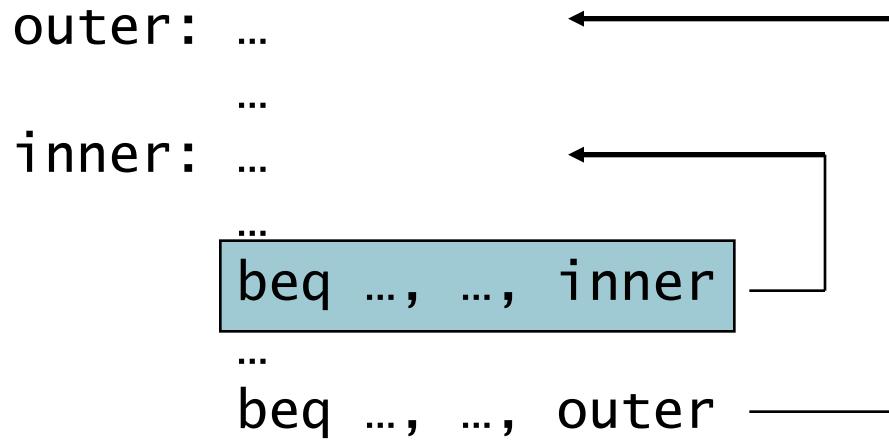
- Static prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Could predict backward branches taken
 - Could predict forward branches not taken
- Dynamic prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - If wrong, take penalty, and update history

Dynamic Branch Prediction

- Dynamic prediction is better
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through (next) or target address
 - If wrong, flush pipeline and flip prediction in the buffer
- In deeper pipelines (more stages), branch penalty is more significant

1-Bit Dynamic Predictor

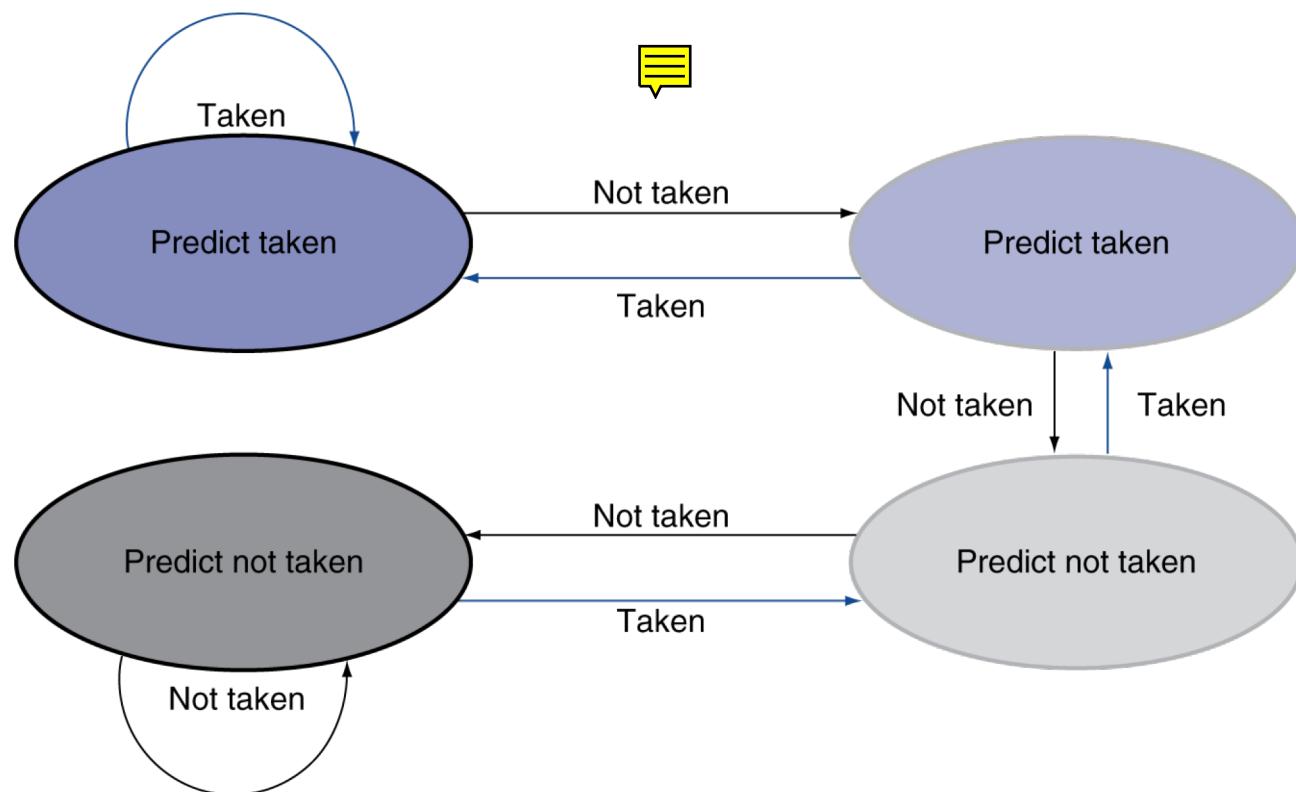
- Assume branch taken
- Inner loop branches mis-predicted twice in every outer loop iteration



- Mis-predict as taken on last iteration of inner loop
- Then mis-predict as not taken on first iteration of inner loop next time around

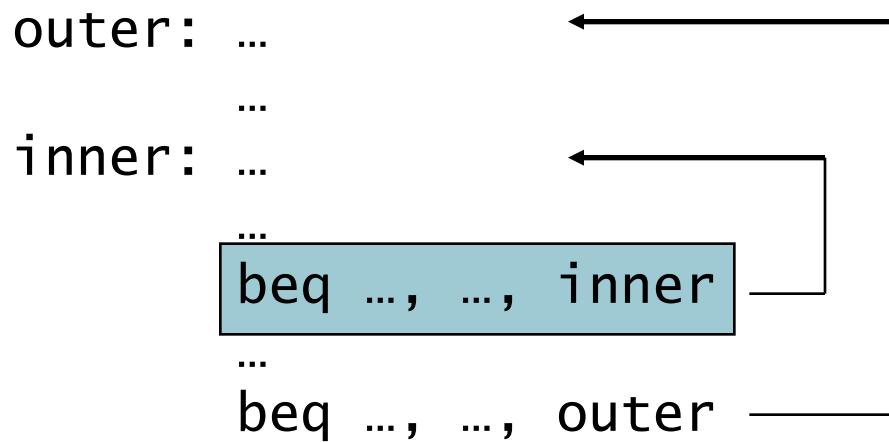
2-Bit Dynamic Predictor

- Only change prediction on two successive mispredictions



2-Bit Dynamic Predictor

- Assume initial state of strong branch taken
- Inner loop branches mis-predicted once in every outer loop iteration



- Mis-predict as taken on last iteration of inner loop

Branch Hazard Resolutions

- Stall on branch
- Always assume branch not taken or taken
- Branch prediction
- *Delayed Branch*

Delayed Branch

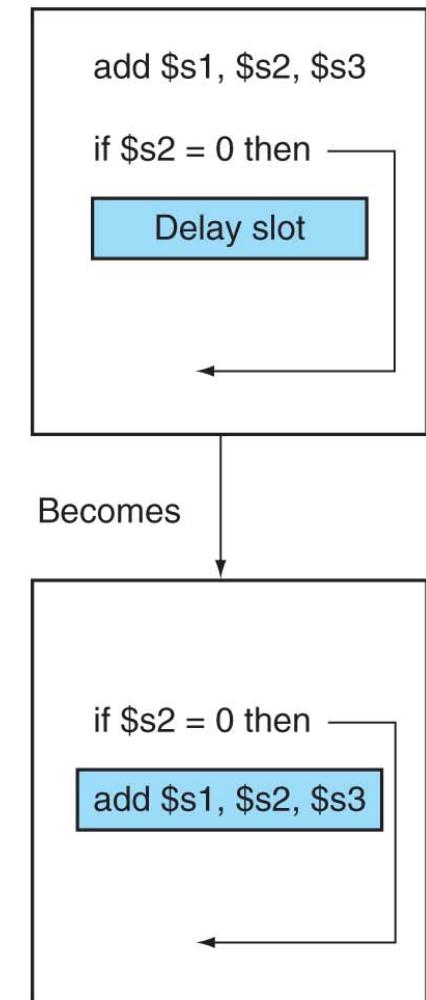
- Always delay the branch
 - With one instruction, NOP or usually a real instruction, for 5 stage pipeline
 - Maybe more delays for deeper pipeline
- Requires carefully designed compiler

Delayed Branch

- Always execute the instruction immediately following branch
 - Called Branch delay slot
 - Further reduce branch penalty
- One branch delay slot becomes insufficient
 - When pipeline gets deeper or multiple instructions per clock cycle
 - Dynamic branch prediction is better choice

Example for Delayed Branch

- Will remove the 1 clock cycle penalty
- Will work only if instructions can be found to fill the delay slot



Overall Structure

